



MACHINE LEARNING

LAB PROJECT SUBMISSION

SUBMITTED TO : Dr. Simran Setia

SUBMITTED BY :

NAME	ROLL NO
Pranav Sharma	102116090

INDEX

INDEX OF THE PROJECT

CONTENT	PAGE NO.
Introduction	2
base_model.py	3
colorization_model.py	10
cycle_gan_model	12
Results	17
Conclusion	18

INTRODUCTION

Classical Art Converter

OVERVIEW

This project employs CycleGAN, a deep learning model, to transform images into the stylistic representation of a specific artist. The objective is to mimic the artistic style of a chosen painter and apply it to arbitrary input images. By leveraging the power of generative adversarial networks (GANs) and cycle consistency, the model aims to learn the nuances of the artist's unique visual language and reproduce them faithfully in the generated outputs.

OBJECTIVE

The primary goal of this project is to enable users to seamlessly translate any image into the distinctive artistic style of a chosen painter. By harnessing the capabilities of CycleGAN, the model seeks to capture the essence of the artist's brushstrokes, color palette, and overall aesthetic, providing a creative tool for users to explore and experiment with diverse artistic expressions.

APPROACH

The project employs the CycleGAN architecture, consisting of two generators and discriminators, to learn the mapping between input images and the target artist's style. The cycle-consistency loss ensures that the transformation is reversible, maintaining content integrity. Training involves optimizing these networks on a curated dataset, emphasizing the importance of both the source and target artistic styles in achieving accurate and visually appealing results.

KEY COMPONENTS:

Data Collection: The dataset comprises a collection of images representing the artistic style of the chosen painter, curated from sources such as online repositories. This dataset enables the model to learn the intricate details of the artist's unique style, facilitating a more robust and generalized transformation process.

Training and Evaluation: Rigorous model training using labeled datasets, followed by meticulous evaluation to measure accuracy, precision, and recall, ensuring the model's efficacy in correctly identifying various yoga postures.

base_model.py

```

1  import os
2  import torch
3  from collections import OrderedDict
4  from abc import ABC, abstractmethod
5  from . import networks
6
7
8  class BaseModel(ABC):
9
10     def __init__(self, opt):
11
12         self.opt = opt
13         self.gpu_ids = opt.gpu_ids
14         self.isTrain = opt.isTrain
15         self.device = torch.device('cuda:{}'.format(self.gpu_ids[0])) if self.gpu_ids else torch.device('cpu') # get device name: CPU or
16         self.save_dir = os.path.join(opt.checkpoints_dir, opt.name) # save all the checkpoints to save_dir
17         if opt.preprocess != 'scale_width': # with [scale_width], input images might have different sizes, which hurts the performance o
18             torch.backends.cudnn.benchmark = True
19         self.loss_names = []
20         self.model_names = []
21         self.visual_names = []
22         self.optimizers = []
23         self.image_paths = []
24         self.metric = 0 # used for learning rate policy 'plateau'
25
26     @staticmethod
27     def modify_commandline_options(parser, is_train):

```

class BaseModel:

This class is an abstract base class (ABC) for models.

To create a subclass, you need to implement the following five functions:

-- <__init__>:	initialize the class; first call BaseModel.__init__(self, opt).
-- <set_input>:	unpack data from dataset and apply preprocessing.
-- <forward>:	produce intermediate results.
-- <optimize_parameters>:	calculate losses, gradients, and update network weights.
-- <modify_commandline_options>:	(optionally) add model-specific options and set default options.

__init__ function:

Initialize the BaseModel class.

Parameters:

opt (Option class)-- stores all the experiment flags; needs to be a subclass of BaseOptions

When creating your custom class, you need to implement your own initialization.

In this function, you should first call <BaseModel.__init__(self, opt)>

Then, you need to define four lists:

-- self.loss_names (str list): specify the training losses that you want to plot and save.

```

        -- self.model_names (str list):          define networks used in our
training.
        -- self.visual_names (str list):         specify the images that you
want to display and save.
        -- self.optimizers (optimizer list):     define and initialize
optimizers. You can define one optimizer for each network. If two networks are
updated at the same time, you can use itertools.chain to group them. See
cycle_gan_model.py for an example.

```

```

46     @staticmethod
47     def modify_commandline_options(parser, is_train):
48
49         return parser
50
51     @abstractmethod
52     def set_input(self, input):
53
54         pass
55
56     @abstractmethod
57     def forward(self):
58         pass
59
60     @abstractmethod
61     def optimize_parameters(self):
62         pass
63
64     def setup(self, opt):
65
66         if self.isTrain:
67             self.schedulers = [networks.get_scheduler(optimizer, opt) for optimizer in self.optimizers]
68         if not self.isTrain or opt.continue_train:
69             load_suffix = 'iter_%d' % opt.load_iter if opt.load_iter > 0 else opt.epoch
70             self.load_networks(load_suffix)
71             self.print_networks(opt.verbose)
72
73     def eval(self):
74         for name in self.model_names:
75             if isinstance(name, str):
76                 net = getattr(self, 'net' + name)
77                 net.eval()
78
79     def test(self):
80
81         with torch.no_grad():
82             self.forward()

```

Modify_commandline_options function:

Add new model-specific options, and rewrite default values for existing options.

Parameters:

```

        parser          -- original option parser
        is_train (bool) -- whether training phase or test phase. You can use
this flag to add training-specific or test-specific options.

```

Returns:

```

        the modified parser.

```

Set_input function:

Unpack input data from the dataloader and perform necessary pre-processing steps.

Parameters:

input (dict): includes the data itself and its metadata information.

forward function:

Run forward pass; called by both functions <optimize_parameters> and <test>.

setup function:

Load and print networks; create schedulers

Parameters:

opt (Option class) -- stores all the experiment flags; needs to be a subclass of BaseOptions

eval function:

Make models eval mode during test time

test function:

Forward function used in test time.

This function wraps <forward> function in no_grad() so we don't save intermediate steps for backprop

It also calls <compute_visuals> to produce additional visualization results

```
98     def test(self):
99
100         with torch.no_grad():
101             self.forward()
102             self.compute_visuals()
103
104     def compute_visuals(self):
105         pass
106
107     def get_image_paths(self):
108         return self.image_paths
109
110     def update_learning_rate(self):
111         old_lr = self.optimizers[0].param_groups[0]['lr']
112         for scheduler in self.schedulers:
113             if self.opt.lr_policy == 'plateau':
114                 scheduler.step(self.metric)
115             else:
116                 scheduler.step()
117
118         lr = self.optimizers[0].param_groups[0]['lr']
119         print('learning rate %.7f -> %.7f' % (old_lr, lr))
120
121     def get_current_visuals(self):
122         visual_ret = OrderedDict()
123         for name in self.visual_names:
124             if isinstance(name, str):
125                 visual_ret[name] = getattr(self, name)
126         return visual_ret
127
128     def get_current_losses(self):
129         errors_ret = OrderedDict()
130         for name in self.loss_names:
131             if isinstance(name, str):
132                 errors_ret[name] = float(getattr(self, 'loss_' + name)) # float(...) works for both scalar tensor and float number
133         return errors_ret
```

compute_visuals function:

Calculate additional output images for visdom and HTML visualization

get_image_paths function:

Return image paths that are used to load current data

update_learning_rate function:

Update learning rates for all the networks; called at the end of every epoch

get_current_visuals function:

Return visualization images. train.py will display these images with visdom, and save the images to a HTML

get_current_losses function:

Return training losses / errors. train.py will print out these errors on console, and save them to a file

```

138     errors_ret = OrderedDict()
139     for name in self.loss_names:
140         if isinstance(name, str):
141             errors_ret[name] = float(getattr(self, 'loss_' + name)) # float(...) works for b
142     return errors_ret
143
144     def save_networks(self, epoch):
145
146         for name in self.model_names:
147             if isinstance(name, str):
148                 save_filename = '%s_net_%s.pth' % (epoch, name)
149                 save_path = os.path.join(self.save_dir, save_filename)
150                 net = getattr(self, 'net' + name)
151
152                 if len(self.gpu_ids) > 0 and torch.cuda.is_available():
153                     torch.save(net.module.cpu().state_dict(), save_path)
154                     net.cuda(self.gpu_ids[0])
155                 else:
156                     torch.save(net.cpu().state_dict(), save_path)
157
158     def __patch_instance_norm_state_dict(self, state_dict, module, keys, i=0):
159         key = keys[i]
160         if i + 1 == len(keys): # at the end, pointing to a parameter/buffer
161             if module.__class__.__name__.startswith('InstanceNorm') and \
162                 (key == 'running_mean' or key == 'running_var'):
163                 if getattr(module, key) is None:
164                     state_dict.pop('.'.join(keys))
165             if module.__class__.__name__.startswith('InstanceNorm') and \
166                 (key == 'num_batches_tracked'):
167                 state_dict.pop('.'.join(keys))
168         else:
169             self.__patch_instance_norm_state_dict(state_dict, getattr(module, key), keys, i + 1)
170
171     def load_networks(self, epoch):

```

save_networks function:

Save all the networks to the disk.

Parameters:

epoch (int) -- current epoch; used in the file name '%s_net_%s.pth' % (epoch, name)

__patch_instance_norm_state_dict function:

Fix InstanceNorm checkpoints incompatibility (prior to 0.4)

```

173         else:
174             self.__patch_instance_norm_state_dict(state_dict, getattr(module, key), keys, i + 1)
175
176     def load_networks(self, epoch):
177
178         for name in self.model_names:
179             if isinstance(name, str):
180                 load_filename = '%s_net_%s.pth' % (epoch, name)
181                 load_path = os.path.join(self.save_dir, load_filename)
182                 net = getattr(self, 'net' + name)
183                 if isinstance(net, torch.nn.DataParallel):
184                     net = net.module
185                 print('loading the model from %s' % load_path)
186                 # if you are using PyTorch newer than 0.4 (e.g., built from
187                 # GitHub source), you can remove str() on self.device
188                 state_dict = torch.load(load_path, map_location=str(self.device))
189                 if hasattr(state_dict, '_metadata'):
190                     del state_dict._metadata
191
192                 # patch InstanceNorm checkpoints prior to 0.4
193                 for key in list(state_dict.keys()): # need to copy keys here because we mutate in loop
194                     self.__patch_instance_norm_state_dict(state_dict, net, key.split('.'))
195                 net.load_state_dict(state_dict)
196
197     def print_networks(self, verbose):
198
199         print('----- Networks initialized -----')
200         for name in self.model_names:
201             if isinstance(name, str):
202                 net = getattr(self, 'net' + name)
203                 num_params = 0
204                 for param in net.parameters():
205                     num_params += param.numel()
206                 if verbose:
207                     print(net)
208                 print('[Network %s] Total number of parameters : %.3f M' % (name, num_params / 1e6))
209         print('-----')

```

load_network function:

Load all the networks from the disk.

Parameters:

epoch (int) -- current epoch; used in the file name '%s_net_%s.pth' % (epoch, name)

print_network function:

Print the total number of parameters in the network and (if verbose) network architecture

Parameters:

verbose (bool) -- if verbose: print the network architecture


```
216         print('[Network %s] Total number of parameters : %.3f M' % (name, num_params / 1e6))
217     print('-----')
218
219     def set_requires_grad(self, nets, requires_grad=False):
220
221         if not isinstance(nets, list):
222             nets = [nets]
223         for net in nets:
224             if net is not None:
225                 for param in net.parameters():
226                     param.requires_grad = requires_grad
```

set_requires_grad function:

Set requires_grad=False for all the networks to avoid unnecessary computations

Parameters:

nets (network list) -- a list of networks

requires_grad (bool) -- whether the networks require gradients or

not

colorization_model.py

```

1  from .pix2pix_model import Pix2PixModel
2  import torch
3  from skimage import color # used for lab2rgb
4  import numpy as np
5
6
7  class ColorizationModel(Pix2PixModel):
8
9      @staticmethod
10     def modify_commandline_options(parser, is_train=True):
11
12         Pix2PixModel.modify_commandline_options(parser, is_train)
13         parser.set_defaults(dataset_mode='colorization')
14         return parser
15
16     def __init__(self, opt):
17
18         # reuse the pix2pix model
19         Pix2PixModel.__init__(self, opt)
20         # specify the images to be visualized.
21         self.visual_names = ['real_A', 'real_B_rgb', 'fake_B_rgb']
22
23     def lab2rgb(self, L, AB):
24
25         AB2 = AB * 110.0
26         L2 = (L + 1.0) * 50.0
27         Lab = torch.cat([L2, AB2], dim=1)
28         Lab = Lab[0].data.cpu().float().numpy()
29         Lab = np.transpose(Lab.astype(np.float64), (1, 2, 0))
30         rgb = color.lab2rgb(Lab) * 255
31         return rgb
32
33     def compute_visuals(self):
34         self.real_B_rgb = self.lab2rgb(self.real_A, self.real_B)
35         self.fake_B_rgb = self.lab2rgb(self.real_A, self.fake_B)
36

```

class ColorizationModel:

This is a subclass of Pix2PixModel for image colorization (black & white image -> colorful images).

The model training requires '-dataset_model colorization' dataset.

It trains a pix2pix model, mapping from L channel to ab channels in Lab color space.

By default, the colorization dataset will automatically set '--input_nc 1' and '--output_nc 2'.

modify_commandline_options function:

Add new dataset-specific options, and rewrite default values for existing options.

Parameters:

parser -- original option parser

is_train (bool) -- whether training phase or test phase. You can use this flag to add training-specific or test-specific options.

Returns:

the modified parser.

By default, we use 'colorization' dataset for this model.

See the original pix2pix paper (<https://arxiv.org/pdf/1611.07004.pdf>) and colorization results (Figure 9 in the paper)

__init__ function:

Initialize the class.

Parameters:

opt (Option class)-- stores all the experiment flags; needs to be a subclass of BaseOptions

For visualization, we set 'visual_names' as 'real_A' (input real image), 'real_B_rgb' (ground truth RGB image), and 'fake_B_rgb' (predicted RGB image)

We convert the Lab image 'real_B' (inherited from Pix2pixModel) to a RGB image 'real_B_rgb'.

we convert the Lab image 'fake_B' (inherited from Pix2pixModel) to a RGB image 'fake_B_rgb'.

lab2rgb function:

Convert an Lab tensor image to a RGB numpy output

Parameters:

L (1-channel tensor array): L channel images (range: [-1, 1], torch tensor array)

AB (2-channel tensor array): ab channel images (range: [-1, 1], torch tensor array)

Returns:

rgb (RGB numpy image): rgb output images (range: [0, 255], numpy array)

compute_visuals function:

Calculate additional output images for visdom and HTML visualization

cycle_gan_model.py

```

1  import torch
2  import itertools
3  from util.image_pool import ImagePool
4  from .base_model import BaseModel
5  from . import networks
6
7
8  class CycleGANModel(BaseModel):
9
10     @staticmethod
11     def modify_commandline_options(parser, is_train=True):
12
13         parser.set_defaults(no_dropout=True) # default CycleGAN did not use dropout
14         if is_train:
15             parser.add_argument('--lambda_A', type=float, default=10.0, help='weight for cycle loss (A -> B -> A)')
16             parser.add_argument('--lambda_B', type=float, default=10.0, help='weight for cycle loss (B -> A -> B)')
17             parser.add_argument('--lambda_identity', type=float, default=0.5, help='use identity mapping. Setting lambda_identity other
18
19         return parser
20
21     def __init__(self, opt):
22
23         BaseModel.__init__(self, opt)
24         # specify the training losses you want to print out. The training/test scripts will call <BaseModel.get_current_losses>
25         self.loss_names = ['D_A', 'G_A', 'cycle_A', 'idt_A', 'D_B', 'G_B', 'cycle_B', 'idt_B']
26         # specify the images you want to save/display. The training/test scripts will call <BaseModel.get_current_visuals>
27         visual_names_A = ['real_A', 'fake_B', 'rec_A']
28         visual_names_B = ['real_B', 'fake_A', 'rec_B']
29         if self.isTrain and self.opt.lambda_identity > 0.0: # if identity loss is used, we also visualize idt_B=G_A(B) ad idt_A=G_A(B)
30             visual_names_A.append('idt_B')
31             visual_names_B.append('idt_A')
32
33         self.visual_names = visual_names_A + visual_names_B # combine visualizations for A and B
34         # specify the models you want to save to the disk. The training/test scripts will call <BaseModel.save_networks> and <BaseModel
35         if self.isTrain:
36             self.model_names = ['G_A', 'G_B', 'D_A', 'D_B']

```

class CycleGANModel:

This class implements the CycleGAN model, for learning image-to-image translation without paired data.

The model training requires '--dataset_mode unaligned' dataset. By default, it uses a '--netG resnet_9blocks' ResNet generator, a '--netD basic' discriminator (PatchGAN introduced by pix2pix), and a least-square GANs objective ('--gan_mode lsgan').

CycleGAN paper: <https://arxiv.org/pdf/1703.10593.pdf>

modify_commandline_options function:

Add new dataset-specific options, and rewrite default values for existing options.

Parameters:

parser -- original option parser

is_train (bool) -- whether training phase or test phase. You can use this flag to add training-specific or test-specific options.

Returns:

the modified parser.

For CycleGAN, in addition to GAN losses, we introduce λ_A , λ_B , and $\lambda_{identity}$ for the following losses.

A (source domain), B (target domain).

Generators: $G_A: A \rightarrow B$; $G_B: B \rightarrow A$.

Discriminators: $D_A: G_A(A)$ vs. B; $D_B: G_B(B)$ vs. A.

Forward cycle loss: $\lambda_A * ||G_B(G_A(A)) - A||$ (Eqn. (2) in the paper)

Backward cycle loss: $\lambda_B * ||G_A(G_B(B)) - B||$ (Eqn. (2) in the paper)

Identity loss (optional): $\lambda_{identity} * (||G_A(B) - B|| * \lambda_B + ||G_B(A) - A|| * \lambda_A)$ (Sec 5.2 "Photo generation from paintings" in the paper)

Dropout is not used in the original CycleGAN paper.

__init__ function:

Initialize the CycleGAN class.

Parameters:

opt (Option class)-- stores all the experiment flags; needs to be a subclass of BaseOptions

```

65     if self.isTrain:
66         self.model_names = ['G_A', 'G_B', 'D_A', 'D_B']
67     else: # during test time, only load Gs
68         self.model_names = ['G_A', 'G_B']
69
70     # define networks (both Generators and discriminators)
71     # The naming is different from those used in the paper.
72     # Code (vs. paper): G_A (G), G_B (F), D_A (D_Y), D_B (D_X)
73     self.netG_A = networks.define_G(opt.input_nc, opt.output_nc, opt.ngf, opt.netG, opt.norm,
74                                     not opt.no_dropout, opt.init_type, opt.init_gain, self.gpu_ids)
75     self.netG_B = networks.define_G(opt.output_nc, opt.input_nc, opt.ngf, opt.netG, opt.norm,
76                                     not opt.no_dropout, opt.init_type, opt.init_gain, self.gpu_ids)
77
78     if self.isTrain: # define discriminators
79         self.netD_A = networks.define_D(opt.output_nc, opt.ndf, opt.netD,
80                                         opt.n_layers_D, opt.norm, opt.init_type, opt.init_gain, self.gpu_ids)
81         self.netD_B = networks.define_D(opt.input_nc, opt.ndf, opt.netD,
82                                         opt.n_layers_D, opt.norm, opt.init_type, opt.init_gain, self.gpu_ids)
83
84     if self.isTrain:
85         if opt.lambda_identity > 0.0: # only works when input and output images have the same number of channels
86             assert(opt.input_nc == opt.output_nc)
87             self.fake_A_pool = ImagePool(opt.pool_size) # create image buffer to store previously generated images
88             self.fake_B_pool = ImagePool(opt.pool_size) # create image buffer to store previously generated images
89             # define loss functions
90             self.criterionGAN = networks.criterionGAN # define GAN loss.
91             self.criterionCycle = torch.nn.L1Loss()
92             self.criterionIdt = torch.nn.L1Loss()
93             # initialize optimizers; schedulers will be automatically created by function <BaseModel.setup>.
94             self.optimizer_G = torch.optim.Adam(itertools.chain(self.netG_A.parameters(), self.netG_B.parameters()), lr=opt.lr, betas=(opt
95             self.optimizer_D = torch.optim.Adam(itertools.chain(self.netD_A.parameters(), self.netD_B.parameters()), lr=opt.lr, betas=(opt
96             self.optimizers.append(self.optimizer_G)
97             self.optimizers.append(self.optimizer_D)
98
99     def set_input(self, input):

```

```

99     def set_input(self, input):
100
101         AtoB = self.opt.direction == 'AtoB'
102         self.real_A = input['A' if AtoB else 'B'].to(self.device)
103         self.real_B = input['B' if AtoB else 'A'].to(self.device)
104         self.image_paths = input['A_paths' if AtoB else 'B_paths']
105
106     def forward(self):
107         self.fake_B = self.netG_A(self.real_A) # G_A(A)
108         self.rec_A = self.netG_B(self.fake_B) # G_B(G_A(A))
109         self.fake_A = self.netG_B(self.real_B) # G_B(B)
110         self.rec_B = self.netG_A(self.fake_A) # G_A(G_B(B))
111
112     def backward_D_basic(self, netD, real, fake):
113
114         # Real
115         pred_real = netD(real)
116         loss_D_real = self.criterionGAN(pred_real, True)
117         # Fake
118         pred_fake = netD(fake.detach())
119         loss_D_fake = self.criterionGAN(pred_fake, False)
120         # Combined loss and calculate gradients
121         loss_D = (loss_D_real + loss_D_fake) * 0.5
122         loss_D.backward()
123         return loss_D
124
125     def backward_D_A(self):
126         fake_B = self.fake_B_pool.query(self.fake_B)
127         self.loss_D_A = self.backward_D_basic(self.netD_A, self.real_B, fake_B)
128
129     def backward_D_B(self):
130         fake_A = self.fake_A_pool.query(self.fake_A)
131         self.loss_D_B = self.backward_D_basic(self.netD_B, self.real_A, fake_A)
132
133     def backward_G(self):
134         lambda_idt = self.opt.lambda_identity
135         lambda_A = self.opt.lambda_A

```

set_input function:

Unpack input data from the dataloader and perform necessary pre-processing steps.

Parameters:

input (dict): include the data itself and its metadata information.

The option 'direction' can be used to swap domain A and domain B.

forward function:

Run forward pass; called by both functions <optimize_parameters> and <test>.

backward_D_basic function:

Calculate GAN loss for the discriminator

Parameters:

```
netD (network)      -- the discriminator D
real (tensor array) -- real images
fake (tensor array) -- images generated by a generator
```

Return the discriminator loss.

We also call `loss_D.backward()` to calculate the gradients.

backward_D_A function:

Calculate GAN loss for discriminator D_A

backward_D_B function:

Calculate GAN loss for discriminator D_B

backward_G function:

Calculate the loss for generators G_A and G_B

```
151 def backward_G(self):
152     lambda_idt = self.opt.lambda_identity
153     lambda_A = self.opt.lambda_A
154     lambda_B = self.opt.lambda_B
155     # Identity loss
156     if lambda_idt > 0:
157         # G_A should be identity if real_B is fed: ||G_A(B) - B||
158         self.idt_A = self.netG_A(self.real_B)
159         self.loss_idt_A = self.criterionIdt(self.idt_A, self.real_B) * lambda_B * lambda_idt
160         # G_B should be identity if real_A is fed: ||G_B(A) - A||
161         self.idt_B = self.netG_B(self.real_A)
162         self.loss_idt_B = self.criterionIdt(self.idt_B, self.real_A) * lambda_A * lambda_idt
163     else:
164         self.loss_idt_A = 0
165         self.loss_idt_B = 0
166
167     # GAN loss D_A(G_A(A))
168     self.loss_G_A = self.criterionGAN(self.netD_A(self.fake_B), True)
169     # GAN loss D_B(G_B(B))
170     self.loss_G_B = self.criterionGAN(self.netD_B(self.fake_A), True)
171     # Forward cycle loss || G_B(G_A(A)) - A ||
172     self.loss_cycle_A = self.criterionCycle(self.rec_A, self.real_A) * lambda_A
173     # Backward cycle loss || G_A(G_B(B)) - B ||
174     self.loss_cycle_B = self.criterionCycle(self.rec_B, self.real_B) * lambda_B
175     # combined loss and calculate gradients
176     self.loss_G = self.loss_G_A + self.loss_G_B + self.loss_cycle_A + self.loss_cycle_B + self.loss_idt_A + self.loss_idt_B
177     self.loss_G.backward()
```

```
180 def optimize_parameters(self):
181     # forward
182     self.forward()      # compute fake images and reconstruction images.
183     # G_A and G_B
184     self.set_requires_grad([self.netD_A, self.netD_B], False) # Ds require no gradients when optimizing Gs
185     self.optimizer_G.zero_grad() # set G_A and G_B's gradients to zero
186     self.backward_G()           # calculate gradients for G_A and G_B
187     self.optimizer_G.step()     # update G_A and G_B's weights
188     # D_A and D_B
189     self.set_requires_grad([self.netD_A, self.netD_B], True)
190     self.optimizer_D.zero_grad() # set D_A and D_B's gradients to zero
191     self.backward_D_A()          # calculate gradients for D_A
192     self.backward_D_B()          # calculate gradients for D_B
193     self.optimizer_D.step()     # update D_A and D_B's weights
194
```

optimize_parameters function:

Calculate losses, gradients, and update network weights; called in every training iteration

RESULTS

RESULTS OF THE PROJECT

MODEL PERFORMANCE AND TRAINING

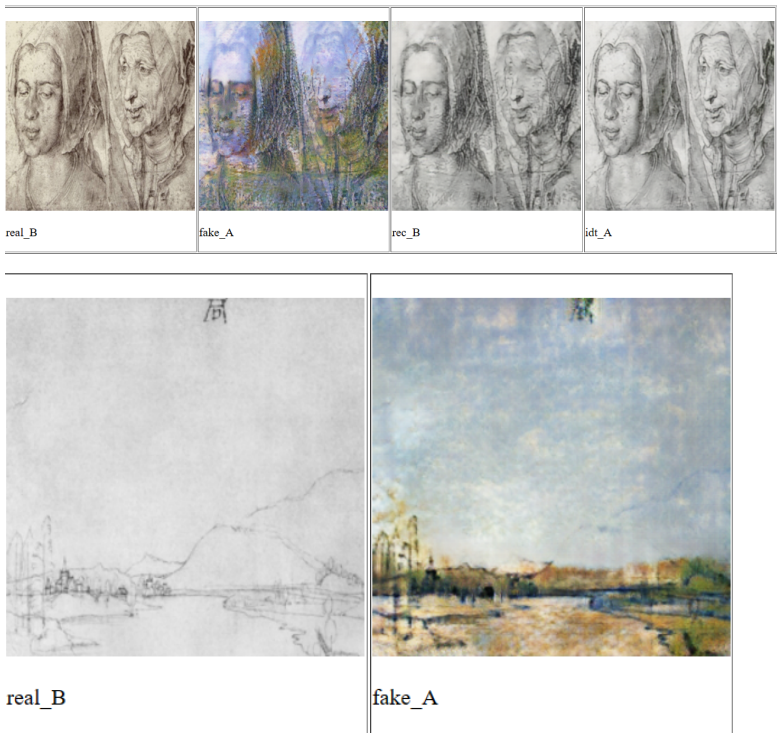
Visual Comparison:

Generated vs. Ground Truth:

epoch [174]



Domain-Specific Evaluation: Edges and Color Correction:



Highlights the effectiveness of discriminators specialized in evaluating edges and color correctness.
Provide examples showcasing enhanced details and color accuracy.

Style Limitations: Classical art generators using CycleGANs may struggle to handle diverse or abstract artistic styles, leading to challenges in faithfully reproducing certain genres.

Detail Loss: Fine details in high-resolution artworks may be lost during the image translation process, affecting the overall fidelity of the generated images.

CONCLUSION

In conclusion, this project successfully leverages CycleGAN for artistic style transfer, providing a valuable tool for users to explore and apply the distinctive styles of renowned painters to their own images. The project contributes to the intersection of deep learning and creative expression, opening avenues for further research and applications in the domain of computer-assisted art creation.